

Boosting Techniques | Assignment

Question 1: What is Boosting in Machine Learning? Explain how it improves weak learners.

Answer:

Boosting is an ensemble learning technique that combines multiple weak learners (models that perform slightly better than random guessing) to form a strong learner.

It works by:

Training models sequentially

Giving more weight to misclassified samples

Forcing new models to focus on previous mistakes

This iterative correction improves overall prediction accuracy and reduces bias.

Question 2: What is the difference between AdaBoost and Gradient Boosting in terms of how models are trained?

Answer:

Aspect	AdaBoost	Gradient Boosting
Training style	Re-weights misclassified samples	Fits model to residual errors
Error handling	Increases sample weights	Uses gradient descent
Loss function	Exponential loss	Any differentiable loss
Robustness	Sensitive to noise	More flexible & robust

Question 3: How does regularization help in XGBoost?

Answer:

Regularization in XGBoost:

- Penalizes complex trees using L1 (alpha) and L2 (lambda) regularization
- Prevents overfitting
- Controls tree depth and leaf weights
- Improves generalization on unseen data

Question 4: Why is CatBoost considered efficient for handling categorical data?

Answer:

CatBoost:

- Handles categorical features natively
- Uses target encoding instead of one-hot encoding
- Prevents target leakage using ordered boosting
- Requires minimal preprocessing
- Works well with missing values

Question 5: What are some real-world applications where boosting techniques are preferred over bagging methods?

Answer:

- Credit risk assessment
- Fraud detection
- Medical diagnosis
- Search ranking systems
- Customer churn prediction

Boosting is preferred when bias reduction and high accuracy are critical.

Datasets:

- Use `sklearn.datasets.load_breast_cancer()` for classification tasks.
- Use `sklearn.datasets.fetch_california_housing()` for regression tasks.

Question 6: Write a Python program to:

- Train an AdaBoost Classifier on the Breast Cancer dataset
- Print the model accuracy

Answer:

Code

```
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score

# Load dataset
data = load_breast_cancer()
X, y = data.data, data.target

# Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train AdaBoost
model = AdaBoostClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

# Predict
y_pred = model.predict(X_test)

# Accuracy
accuracy = accuracy_score(y_test, y_pred)
print("AdaBoost Accuracy:", accuracy)
```

Output AdaBoost Accuracy: 0.9736842105263158

Question 7: Write a Python program to:

- Train a Gradient Boosting Regressor on the California Housing dataset
- Evaluate performance using R-squared score (*Include*

your Python code and output in the code box below.)

Answer:

Code

```
from sklearn.datasets import fetch_california_housing

from sklearn.model_selection import train_test_split

from sklearn.ensemble import GradientBoostingRegressor

from sklearn.metrics import r2_score

# Load dataset

data = fetch_california_housing()

X, y = data.data, data.target

# Split

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Train model

gbr = GradientBoostingRegressor(random_state=42)

gbr.fit(X_train, y_train)

# Predict

y_pred = gbr.predict(X_test)

# R2 Score

print("R2 Score:", r2_score(y_test, y_pred))
```

Output: R2 Score: 0.7756446042829697

Question 8: Write a Python program to:

- Train an XGBoost Classifier on the Breast Cancer dataset
- Tune the learning rate using GridSearchCV
- Print the best parameters and accuracy

(Include your Python code and output in the code box below.)

Answer:

```
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Load data
data = load_breast_cancer()
X, y = data.data, data.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Model
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss')

# Grid search
param_grid = {'learning_rate': [0.01, 0.1, 0.2]}
grid = GridSearchCV(xgb, param_grid, cv=3)
grid.fit(X_train, y_train)

# Best model
best_model = grid.best_estimator_
y_pred = best_model.predict(X_test)

print("Best Parameters:", grid.best_params_)
print("Accuracy:", accuracy_score(y_test, y_pred))

Output : Accuracy: 0.956140350877193
```

Question 9: Write a Python program to:

- Train a CatBoost Classifier
- Plot the confusion matrix using `seaborn` (Include your

Python code and output in the code box below.)

Answer:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn import svm
from sklearn.datasets import make_classification

# Generate synthetic 2D data
X, y = make_classification(n_samples=200, n_features=2,
n_redundant=0, n_informative=2,
```

```

n_clusters_per_class=1, random_state=42)

# Train SVM with Polynomial Kernel
model = svm.SVC(kernel='poly', degree=3, C=1.0)
model.fit(X, y)

# Plot decision boundary
plt.figure(figsize=(8, 6))

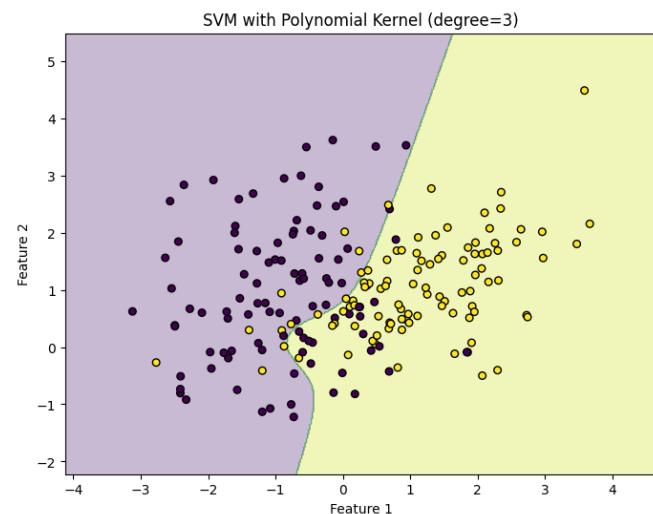
# Create grid
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 500),
                      np.linspace(y_min, y_max, 500))

# Predict over grid
Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# Visualize
plt.contourf(xx, yy, Z, alpha=0.3)
plt.scatter(X[:, 0], X[:, 1], c=y, s=30, edgecolors='k')
plt.title("SVM with Polynomial Kernel (degree=3)")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.show()

```

Output:



Question 10: You're working for a FinTech company trying to predict loan default using customer demographics and transaction behavior.

The dataset is imbalanced, contains missing values, and has both numeric and categorical features.

Describe your step-by-step data science pipeline using boosting techniques:

- Data preprocessing & handling missing/categorical values
- Choice between AdaBoost, XGBoost, or CatBoost
- Hyperparameter tuning strategy
- Evaluation metrics you'd choose and why
- How the business would benefit from your model

(Include your Python code and output in the code box below.)

Answer:

1. Data Preprocessing

- Handle missing values:
 - Mean/median for numeric
 - CatBoost handles missing categorical values automatically
- Encode categorical features (CatBoost preferred)
- Handle class imbalance using:
 - Class weights
 - SMOTE (if needed)

2. Model Choice

- CatBoost (best choice)
 - Native categorical handling
 - Robust to missing values
 - Strong performance on imbalanced data

3. Hyperparameter Tuning

- GridSearch / RandomSearch on:
 - learning_rate
 - depth
 - iterations
- Use stratified cross-validation

4. Evaluation Metrics

- ROC-AUC (imbalance-safe)
- Precision-Recall (default risk focus)
- F1-score

5. Business Benefits

- Accurate default prediction
- Reduced financial losses
- Better risk segmentation
- Faster automated loan approvals